


Training: calibration-aware fine-tuning (CAFT)

every batch imitates the deployment calibration block; all LaBraM layers + head are fine-tuned (same optimiser, schedule and epochs as the baseline)

one batch = one session: $K = 4$ training subjects $\times$ $M = 15$ (clip, second) positions, the same positions for every subject, 3 per emotion — i.e. a calibration block



Deployment on a new user (unchanged)

the same operation the network was trained with; no gradient step at test time

