




(a) Gain needs calibration**(b) 13 of 16 subjects improve****(c) Robust across seeds**

SEED-V, 16 subjects. Standard fine-tuning (std) vs CAFT ① (in-batch centring), 3 seeds each; per-subject values are averaged over seeds.

(a) bars = bootstrap 95% CI of the paired gain; stars = Wilcoxon over subjects (* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$).

(b) full clip; x/16 up = subjects improved; dashed = chance. (c) one point per seed (16-subject mean); CAFT ①+② has seed 0 only.

SLA = stimulus-locked alignment at deployment.