



(a) Cross-subject fine-tuning



(b) Calibration — estimate the domain mean



(c) Inference and scoring — where window and clip differ

evaluation clip (one film clip, N non-overlapping 4 s windows)



N is set by the stimulus (SEED 46–66, SEED-V 13–74, SEED-IV 10–64) — not a choice



one probability vector per window

	neg	neu	pos
win 1	0.09	0.85	0.06
win 2	0.03	0.93	0.04
win 3	0.42	0.49	0.09
⋮	⋮	⋮	⋮
win N	0.03	0.94	0.03

window

window accuracy
each window is one decision — N per clip
long clips therefore weigh more

clip

clip accuracy
mean of the N vectors, then arg max
[0.149, **0.815**, 0.036] → neutral
ONE decision per clip

Averaging the probabilities is not a majority vote: window 3 (0.42 vs 0.49) is nearly undecided and contributes only 0.49, while a confident window contributes 0.94. In this clip 51/58 windows vote neutral (window accuracy 0.879) and the averaged vector also gives neutral — but the two can disagree.

Both are reported throughout. Which one is primary depends on how many evaluation clips a given analysis leaves: clip when there are ≥ 36 , window when only 9 (then one clip is worth 0.111).

— forward pass

- - - - - estimated once per session



proposed: centring $\tilde{z} = z - \mu$