



(a) Leave-one-subject-out



every subject is test exactly once; validation rotates — no RNG
 each fold is run at 3 seeds $\rightarrow 15 \times 3 = 45$ runs
 the epoch is chosen on validation; test selects nothing

(b) Aggregation and testing



(c) Which clips are scored — and that is why the same condition has different values

one session shown; the protocol repeats in all 3 sessions

all 15 clips scored **45 eval clips / subject** *S2, role analysis*

1 clip per emotion held out as calibration **36 eval clips / subject** *ladder, S11, S12*

1 per emotion scored, the rest is the label pool **9 eval clips / subject** *few-shot (S14, S15)*

scored calibration (builds μ) label pool (few-shot only)

A clip that builds μ cannot also be scored.

Primary metric follows this panel: clip when
 ≥ 36 clips are scored, window when only 9
 — one clip would then move it by 0.111.